

# HOST VULNERABILITIES

**UT CS361S – Network Security and Privacy**

**Fall 2021**

Lecture Notes



# VULNERABILITIES AND NETSEC

- This class is “Network Security”
- What do host vulnerabilities have to do with it?
- Hosts are “nodes” in a network graph
- Vulnerabilities can be exploited by remote attackers
  - Either to directly access resources on a particular host
  - Or, to penetrate network defenses and access a more valuable host

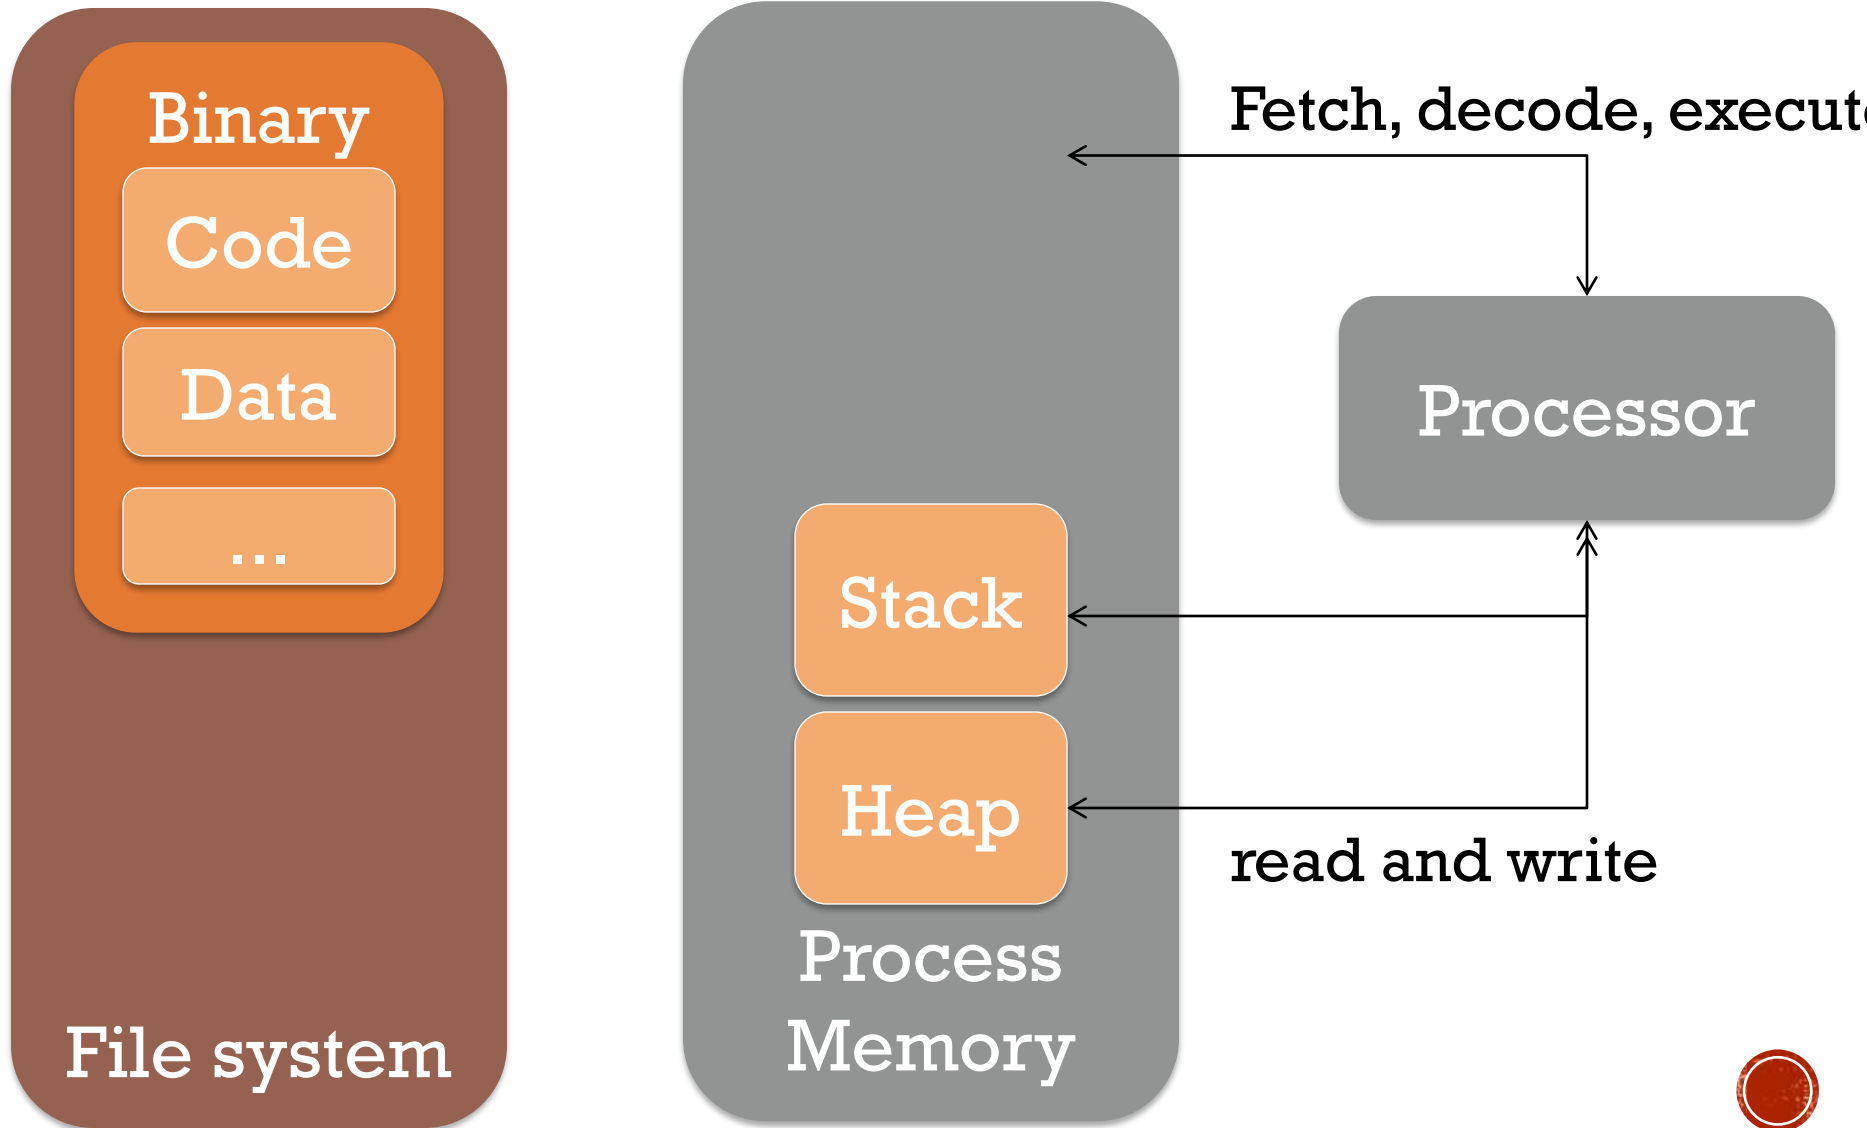


# BRIEF OVERVIEW TO EXECUTION

- ***Very brief*** overview of **Control Flow Hijacking**
  - There are other types of vulnerabilities (e.g misconfigured)
  - Control Flow Hijacking is probably the hardest to grasp
- **Critical Concepts:**
  - The “normal” flow of control for authorized instructions
  - Inputs that change the flow to unauthorized instructions
- **ATtribution:** Derived from slides by Dave Brumley, CMU



# BASIC EXECUTION



# SETH'S NOTES

- Stack

- For temporary static variables
- Function call/return data
- Linear
- Generally, tightly managed

- Heap

- Global variables and dynamic variables
- Hierarchical, “free floating”
- Fragmented, not tightly managed



# SETH'S NOTES

- Assembly “function calls” don’t really exist
  - Rather, jump to new location (“function”)
  - Save context of old location
  - Load context for new location
  - Include information for “returning”

# SETH'S NOTES

- There are multiple ways to do this
- “Calling Conventions”
- Caller Cleanup – caller cleans stack
- Callee Cleanup – called function cleans stack
- Other convention variations:
  - Order that function data is loaded onto stack
  - Whether some data is put into registers instead

# SETH'S NOTES

- Visualizing caller v callee cleanup

```
push arg1
push arg2
push arg3
call proc
```

```
proc:
    pop r1    ; the return address
    pop r2
    pop r2
    pop r2
    push r1
    ret
```

decl (ca

```
push arg1
push arg2
push arg3
call proc
pop r2
pop r2
pop r2
```

```
proc:
    ret
```





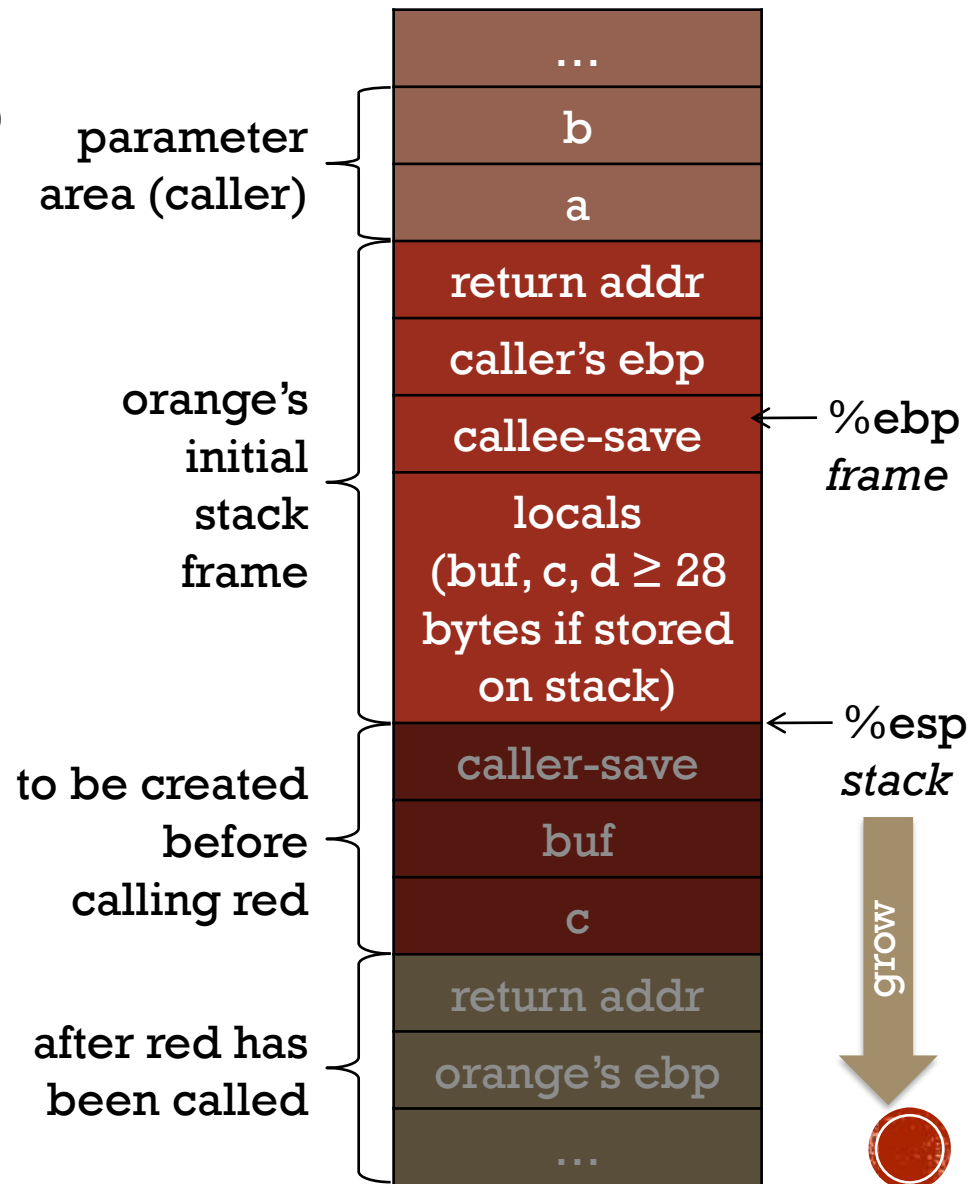
# EBP AND ESP

- **EBP**
  - Stack Base Pointer
  - Where the stack was when the routine started
- **ESP**
  - Stack Pointer
  - Top of the current stack
- **EBP is a previous function's saved ESP**



# CDECL – DEFAULT FOR LINUX & GCC

```
int orange(int a, int b)
{
    char buf[16];
    int c, d;
    if(a > b)
        c = a;
    else
        c = b;
    d = red(c, buf);
    return d;
}
```



# GDB WALKTHROUGH

- SRC:  
[tenouk.com/Bufferoverflowc/Bufferoverflow3.html](http://tenouk.com/Bufferoverflowc/Bufferoverflow3.html)
- Given C code, examine assembly via GDB
- Uses cdecl calling convention



# GDB WALKTHROUGH — C CODE

```
#include <stdio.h>
```

```
int TestFunc(int parameter1, int parameter2, char parameter3)
{
    int y = 3, z = 4;
    char buff[7] = "ABCDEF";

    // function's task code here
    return 0;
}
```

```
int main(int argc, char *argv[ ])
{
    TestFunc(1, 2, 'A');
    return 0;
}
```



# GDB WALKTHROUGH — CALL TESTFUNC

Register `eax` loaded with character 'A' (0x41), not shown

0x08048390 <main+36>: push %eax ;push the third parameter, 'A' prepared in `eax` onto the stack, [ebp+16]

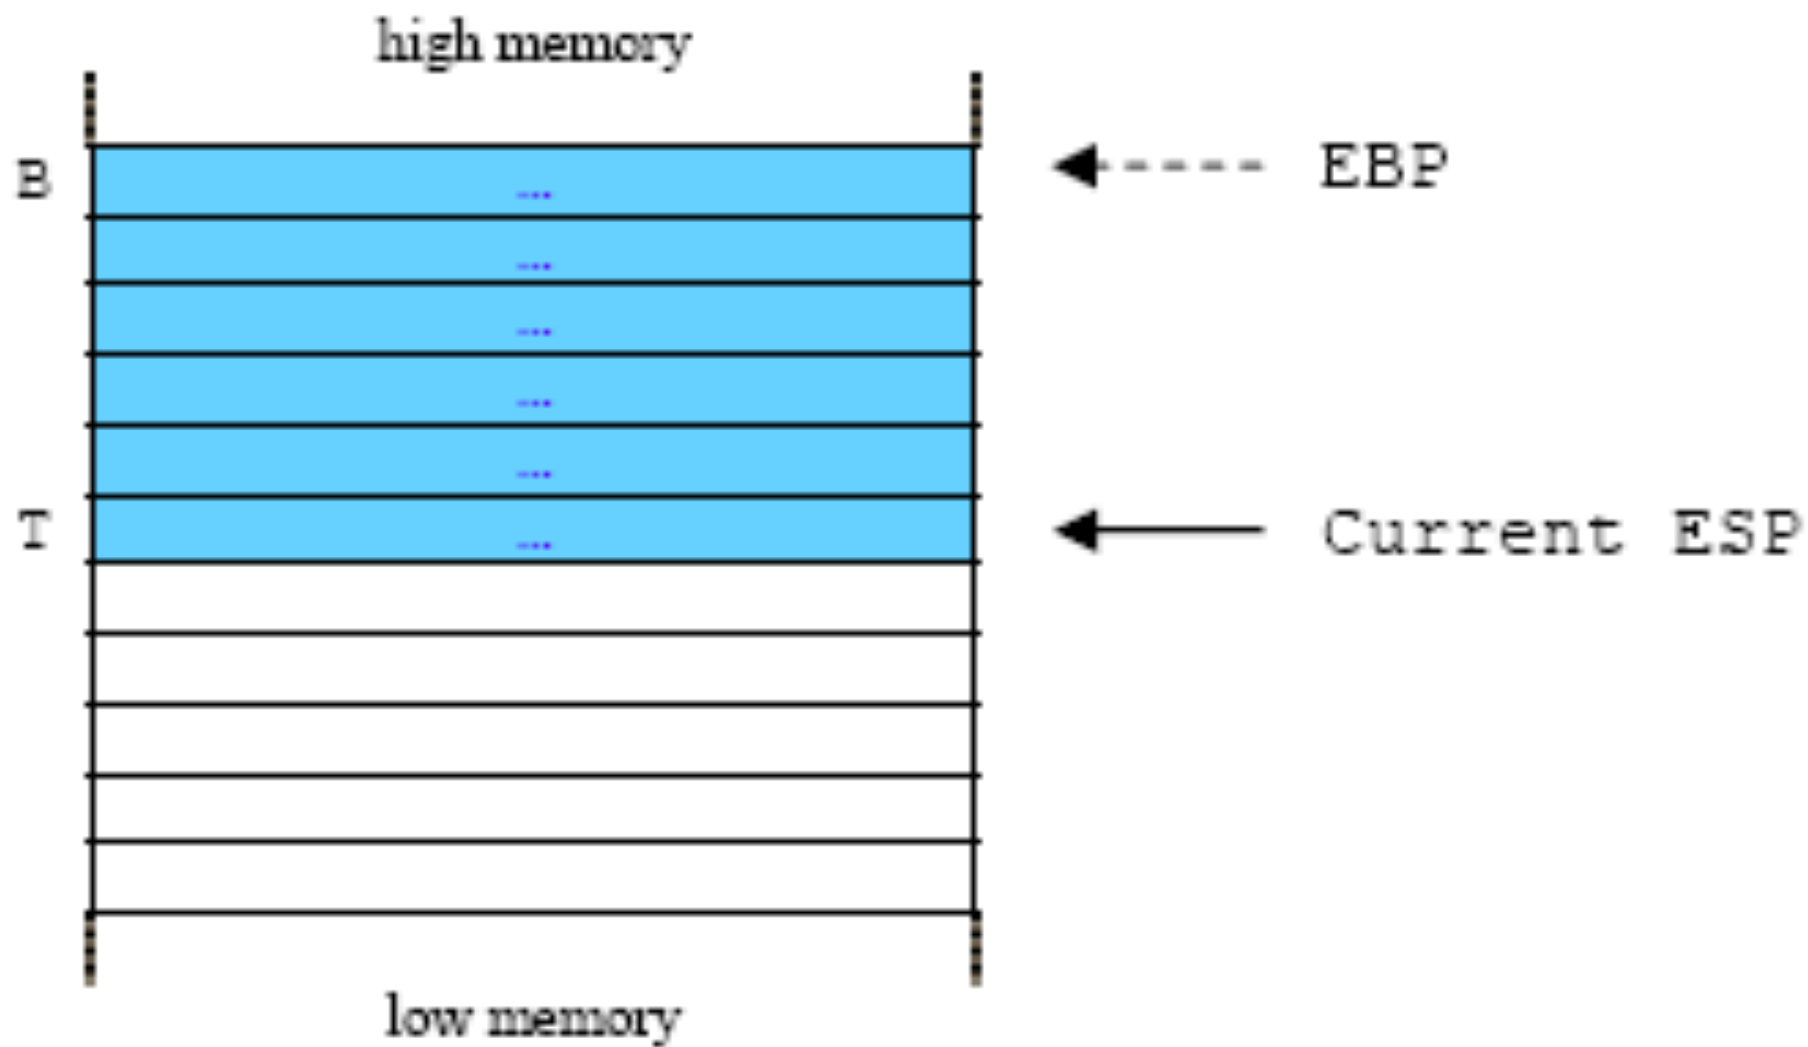
0x08048391 <main+37>: push \$0x2 ;push the second parameter, 2 onto the stack, [ebp+12]


0x08048393 <main+39>: push \$0x1 ;push the first parameter, 1 onto the stack, [ebp+8]

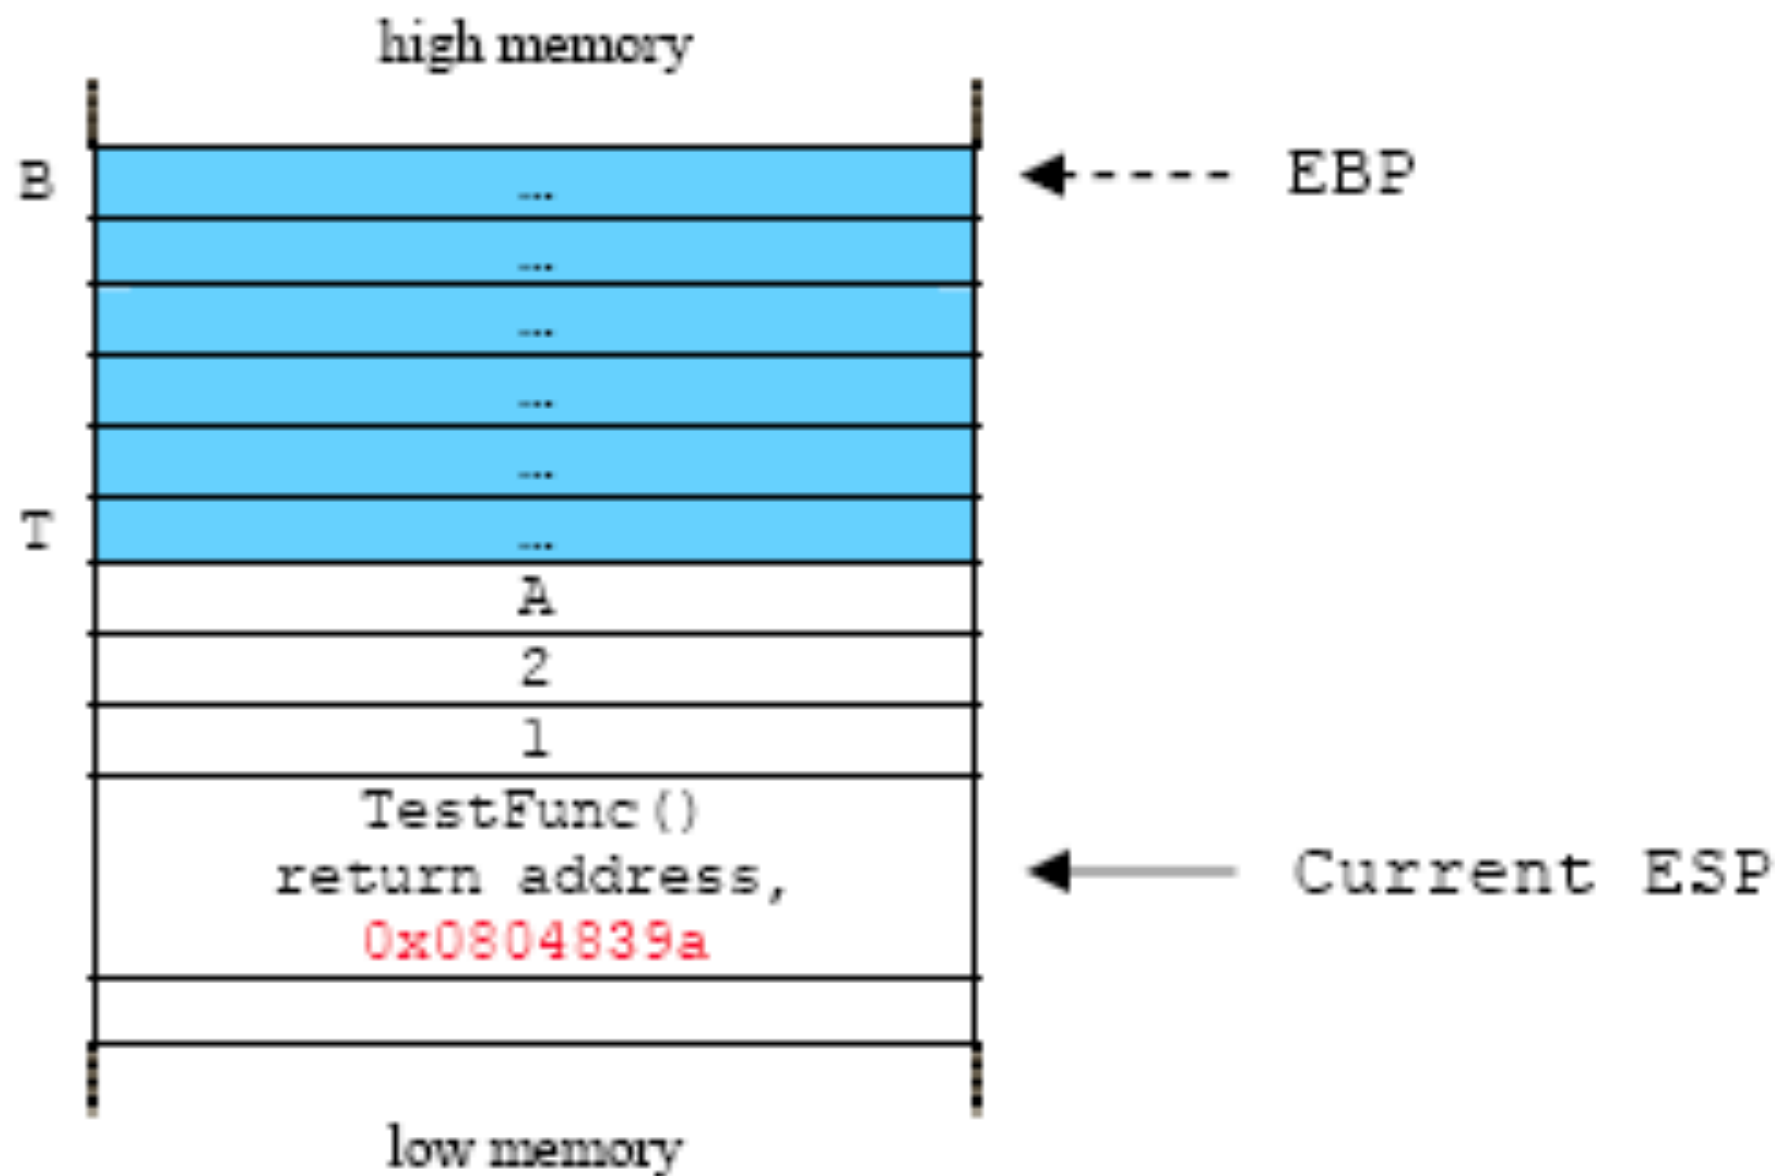
---

0x08048395 <main+41>: call 0x8048334 <TestFunc> ;function call. Push the return  
;address [0x0804839a] onto the stack, [ebp+4]











# GDB WALKTHROUGH – TESTFUNC() C CODE

```
int TestFunc(int parameter1,int parameter2,char parameter3)
{
int y = 3, z = 4;
char buff[7] = "ABCDEF";

// function's task code here
return 0;
}
```

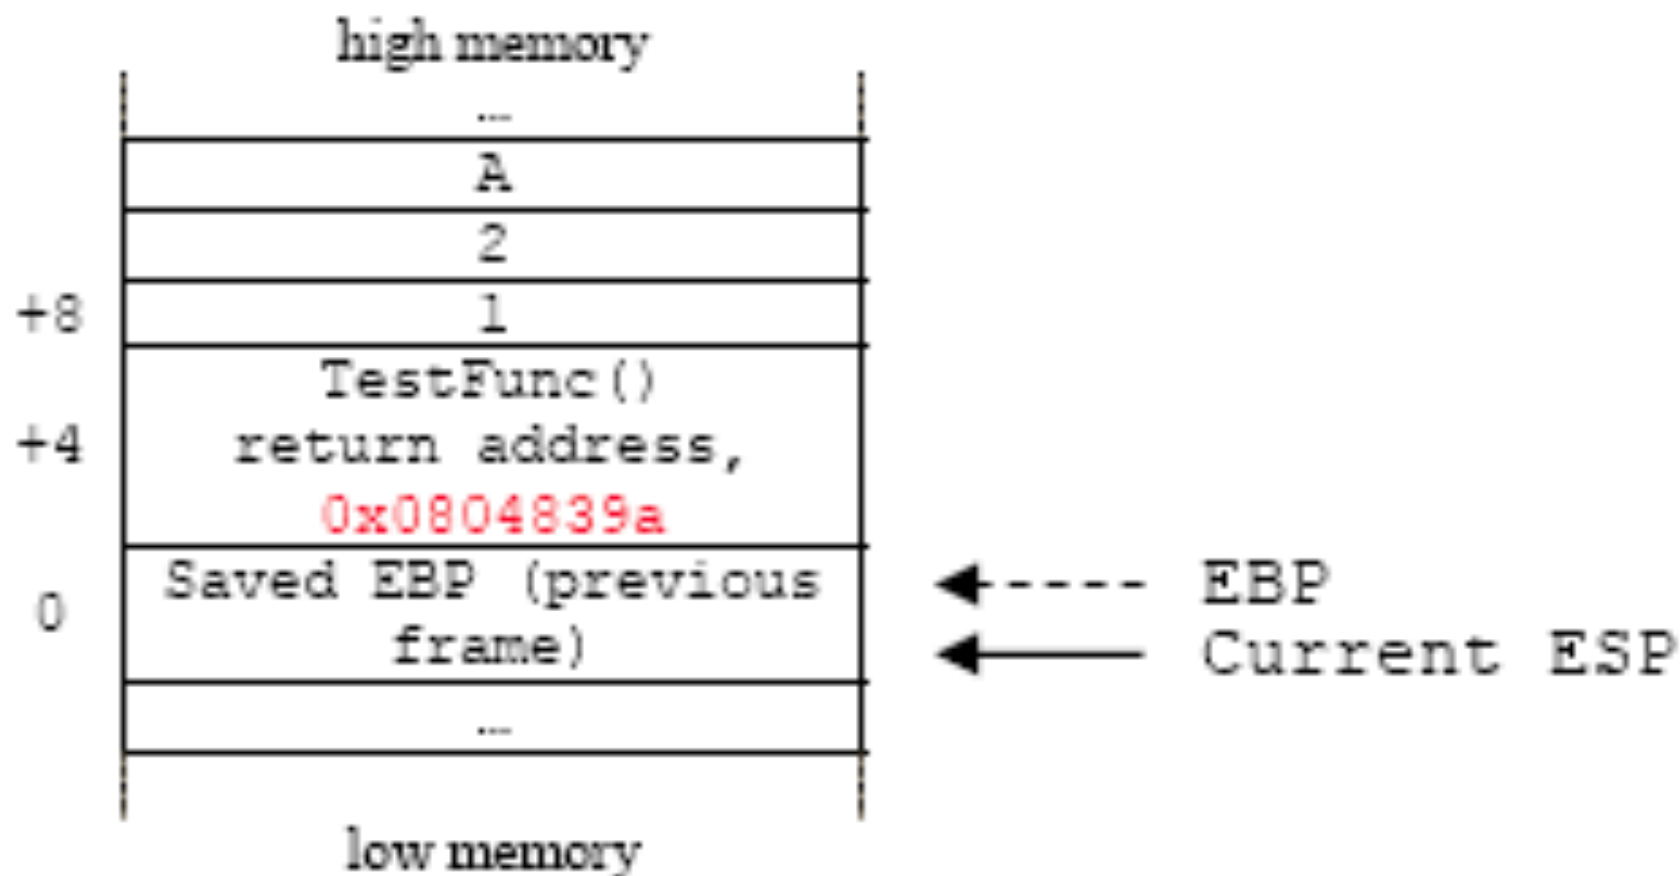


# GDB WALKTHROUGH – TESTFUNC() ASSEMBLY

```
0x08048334 <TestFunc+0>:      push    %ebp                ;push the previous stack frame
                                ;pointer onto the stack, [ebp+0]
0x08048335 <TestFunc+1>:      mov     %esp, %ebp          ;copy the ebp into esp, now the ebp and esp
                                ;are pointing at the same address,
                                ;creating new stack frame [ebp+0]
0x08048337 <TestFunc+3>:      push    %edi                ;save/push edi register, [ebp-4]
0x08048338 <TestFunc+4>:      push    %esi                ;save/push esi register, [ebp-8]
0x08048339 <TestFunc+5>:      sub     $0x20, %esp          ;subtract esp by 32 bytes for local
                                ;variable and buffer if any, go to [ebp-40]
```

32 bytes allocated on stack (0x20). Variables Loaded into this space (not shown).





+4	TestFunc() return address, <b>0x0804839a</b>			
0	Saved EBP (previous frame)			
-4	edi			
-8	esi			
-12	?			
-16	3			
-20	4			
-24				
-28				
-32				
-36	A	B	C	
-40	D	E	F	\0

low memory

←---- EBP

Save Registers

Local variables 'y' and 'z'

buff[7] (note, more allocated)

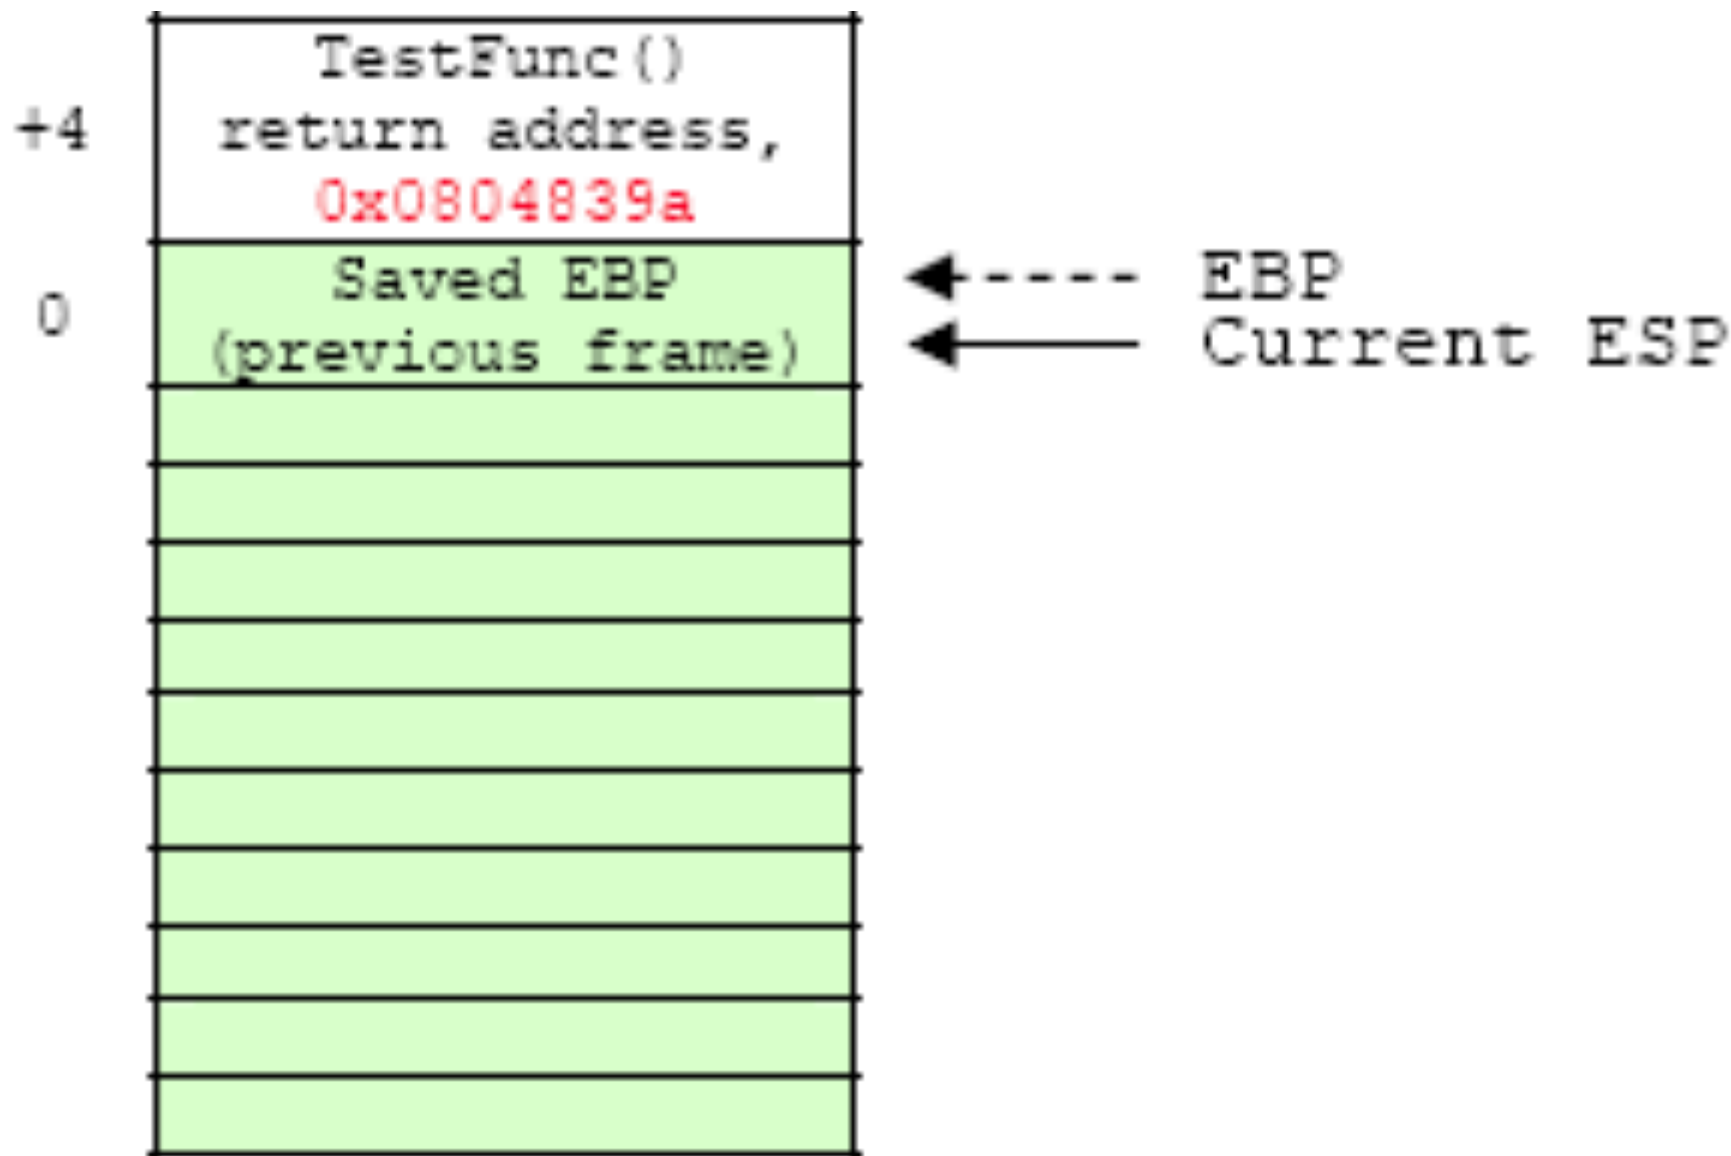
← Current ESP



# GDB WALKTHROUGH — TESTFUNC() EXIT

```
0x08048365 <TestFunc+49>:  add    $0x20, %esp      ;add 32 bytes to esp, back to [ebp-8]
0x08048368 <TestFunc+52>:  pop     %esi            ;restore the esi, [ebp-4]
0x08048369 <TestFunc+53>:  pop     %edi            ;restore the edi, [ebp+0]
```





# GDB WALKTHROUGH – TESTFUNC() EXIT, PART 2

```
0x0804836a <TestFunc+54>:    leave    ;restoring the ebp to the previous stack frame, [ebp+4]
0x0804836b <TestFunc+55>:    ret      ;transfer control back to calling function using
                                ;the saved return address at [ebp+8]
```





# GDB WALTHROUGH — MAIN() AFTER TESTFUNC() RETURN

```
0x0804839a <main+46>:  add    $0xc, %esp    ;cleanup the 3 parameters pushed on the stack  
                        ;at [ebp+8], [ebp+12] and [ebp+16]  
                        ;total up is 12 bytes = 0xc
```







# WHAT ARE BUFFER OVERFLOWS?

A **buffer overflow** occurs when data is written outside of the space allocated for the buffer.

- C does not check that writes are in-bound

## 1. Stack-based

- covered in this class

## 2. Heap-based


- more advanced
- very dependent on system and library version

# BASIC EXAMPLE

```
#include <string.h>
int main(int argc, char **argv) {
    char buf[64];
    strcpy(buf, argv[1]);
}
```

Dump of assembler code for function main:

```
0x080483e4 <+0>:      push    %ebp
0x080483e5 <+1>:      mov     %esp,%ebp
0x080483e7 <+3>:      sub     $72,%esp
0x080483ea <+6>:      mov     12(%ebp),%eax
0x080483ed <+9>:      mov     4(%eax),%eax
0x080483f0 <+12>:     mov     %eax,4(%esp)
0x080483f4 <+16>:     lea     -64(%ebp),%eax
0x080483f7 <+19>:     mov     %eax,(%esp)
0x080483fa <+22>:     call    0x8048300 <strcpy@plt>
0x080483ff <+27>:     leave
0x08048400 <+28>:     ret
```




# "123456"

```
#include <string.h>
int main(int argc, char **argv) {
    char buf[64];
    strcpy(buf, argv[1]);
}
```

Dump of assembler code for function main:

```
0x080483e4 <+0>:      push    %ebp
0x080483e5 <+1>:      mov     %esp,%ebp
0x080483e7 <+3>:      sub     $72,%esp
0x080483ea <+6>:      mov     12(%ebp),%eax
0x080483ed <+9>:      mov     4(%eax),%eax
0x080483f0 <+12>:     mov     %eax,4(%esp)
0x080483f4 <+16>:     lea     -64(%ebp),%eax
0x080483f7 <+19>:     mov     %eax,(%esp)
0x080483fa <+22>:     call    0x8048300 <strcpy@plt>
0x080483ff <+27>:     leave
0x08048400 <+28>:     ret
```

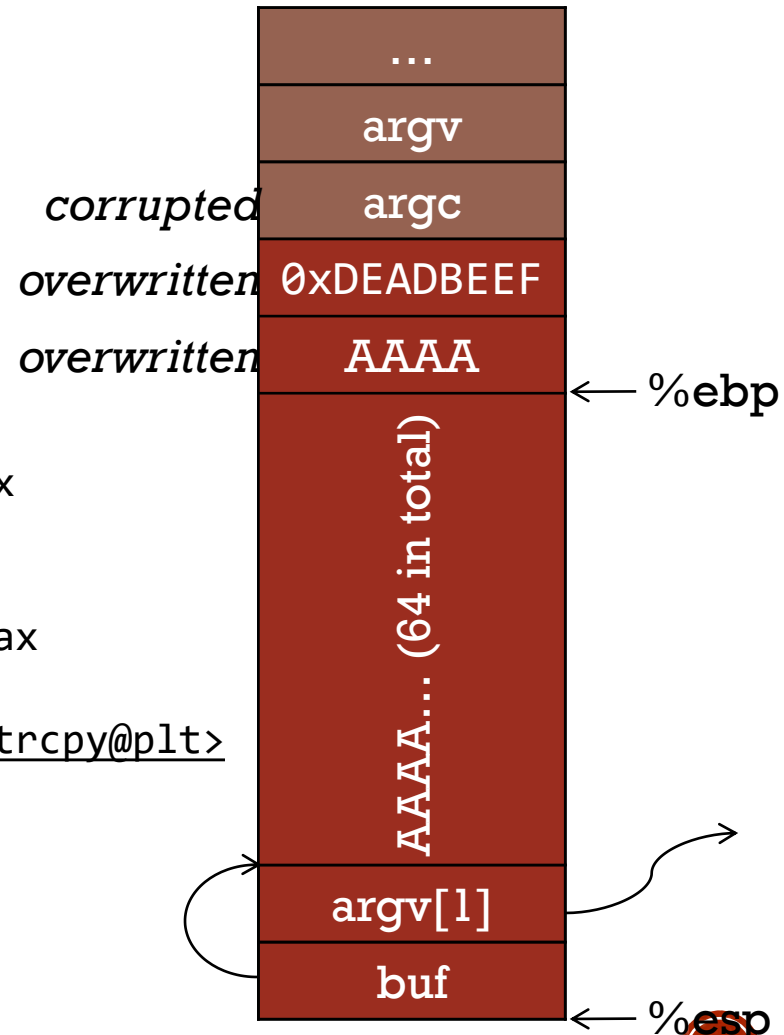


# "A"X68 . "\XEF\XBE\XAD\XDE"

```
#include <string.h>
int main(int argc, char **argv) {
    char buf[64];
    strcpy(buf, argv[1]);
}
```

Dump of assembler code for function main:

```
0x080483e4 <+0>:    push    %ebp
0x080483e5 <+1>:    mov     %esp,%ebp
0x080483e7 <+3>:    sub     $72,%esp
0x080483ea <+6>:    mov     12(%ebp),%eax
0x080483ed <+9>:    mov     4(%eax),%eax
0x080483f0 <+12>:   mov     %eax,4(%esp)
0x080483f4 <+16>:   lea     -64(%ebp),%eax
0x080483f7 <+19>:   mov     %eax,(%esp)
0x080483fa <+22>:   call    0x8048300 <strcpy@plt>
0x080483ff <+27>:   leave
0x08048400 <+28>:   ret
```

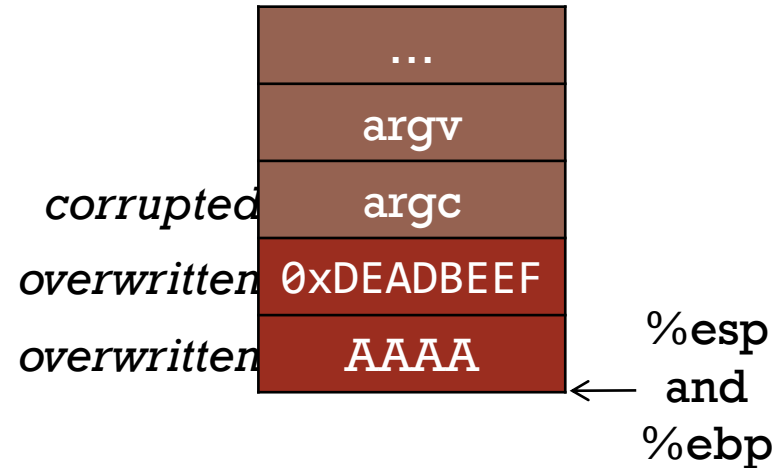


# FRAME TEARDOWN—1

```
#include <string.h>
int main(int argc, char **argv) {
    char buf[64];
    strcpy(buf, argv[1]);
}
```

Dump of assembler code for function main:

```
0x080483e4 <+0>:      push    %ebp
0x080483e5 <+1>:      mov     %esp,%ebp
0x080483e7 <+3>:      sub     $72,%esp
0x080483ea <+6>:      mov     12(%ebp),%eax
0x080483ed <+9>:      mov     4(%eax),%eax
0x080483f0 <+12>:     mov     %eax,4(%esp)
0x080483f4 <+16>:     lea     -64(%ebp),%eax
0x080483f7 <+19>:     mov     %eax,(%esp)
0x080483fa <+22>:     call    0x8048300 <strcpy@plt>
=> 0x080483ff <+27>:     leave
0x08048400 <+28>:     ret
```



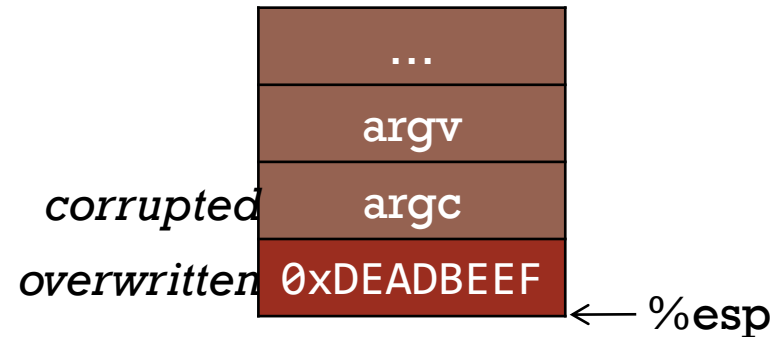
leave  
1. mov %ebp,%esp  
2. pop %ebp

# FRAME TEARDOWN—2

```
#include <string.h>
int main(int argc, char **argv) {
    char buf[64];
    strcpy(buf, argv[1]);
}
```

Dump of assembler code for function main:

```
0x080483e4 <+0>:      push    %ebp
0x080483e5 <+1>:      mov     %esp,%ebp
0x080483e7 <+3>:      sub     $72,%esp
0x080483ea <+6>:      mov     12(%ebp),%eax
0x080483ed <+9>:      mov     4(%eax),%eax
0x080483f0 <+12>:     mov     %eax,4(%esp)
0x080483f4 <+16>:     lea     -64(%ebp),%eax
0x080483f7 <+19>:     mov     %eax,(%esp)
0x080483fa <+22>:     call    0x8048300 <strcpy@plt>
0x080483ff <+27>:     leave
0x08048400 <+28>:     ret
```



%ebp = AAAA


leave  
1. mov %ebp,%esp  
2. pop %ebp

# FRAME TEARDOWN—3

```
#include <string.h>
int main(int argc, char **argv) {
    char buf[64];
    strcpy(buf, argv[1]);
}
```

Dump of assembler code for function main:

```
0x080483e4 <+0>:      push    %ebp
0x080483e5 <+1>:      mov     %esp,%ebp
0x080483e7 <+3>:      sub     $72,%esp
0x080483ea <+6>:      mov     12(%ebp),%eax
0x080483ed <+9>:      mov     4(%eax),%eax
0x080483f0 <+12>:     mov     %eax,4(%esp)
0x080483f4 <+16>:     lea     -64(%ebp),%eax
0x080483f7 <+19>:     mov     %eax,(%esp)
0x080483fa <+22>:     call    0x8048300 <strcpy@plt>
0x080483ff <+27>:     leave
0x08048400 <+28>:     ret
```



*%eip = 0xDEADBEEF  
(probably crash)*



# SHELLCODE

Traditionally, we inject assembly instructions for `exec("/bin/sh")` into buffer.


- see “*Smashing the stack for fun and profit*” for exact string
- or search online

...

```
0x080483fa <+22>: call    0x8048300 <strcpy@plt>
```

```
0x080483ff <+27>: leave
```

```
0x08048400 <+28>: ret
```



# RECAP

To generate ***exploit*** for a basic buffer overflow:

1. Determine size of **stack frame up to head of buffer**
2. Overflow buffer with the right size



***computation***

+

***control***

## DEALING WITH CONTROL FLOW VIOLATIONS



Make it harder to  
control a subverted  
flow



Make taking control of  
the flow innocuous




Make it harder to get  
control of the flow



# DISRUPTING EXPLOITATIVE OPERATIONS



# RANDOM STACK GAP



# ASLR

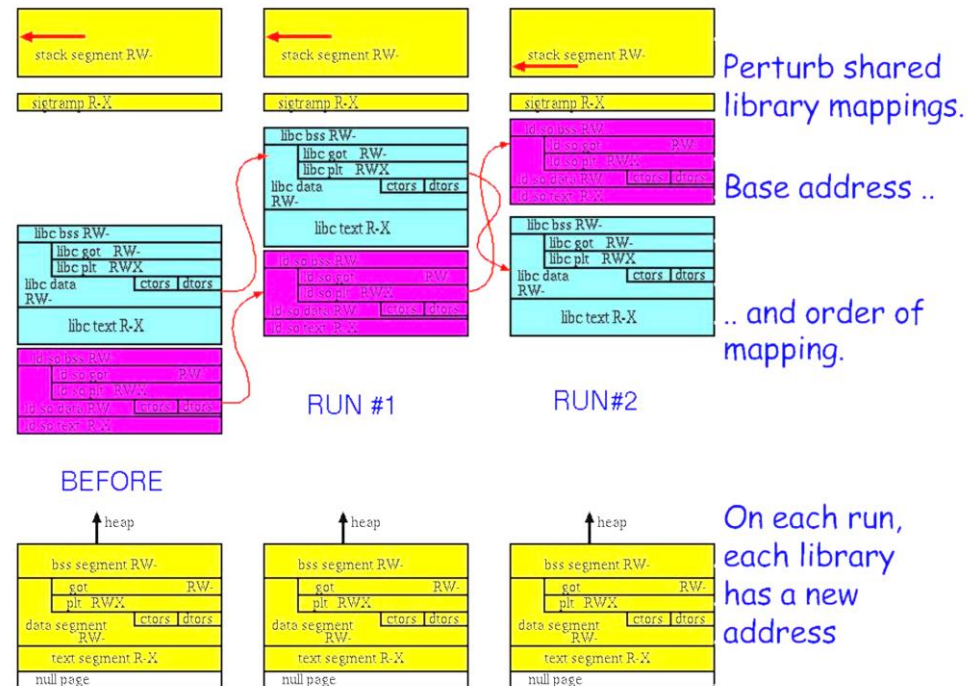
## Address Space Layout Randomization

- Subversion usually needs to know memory layout
- General goal: make layout unpredictable

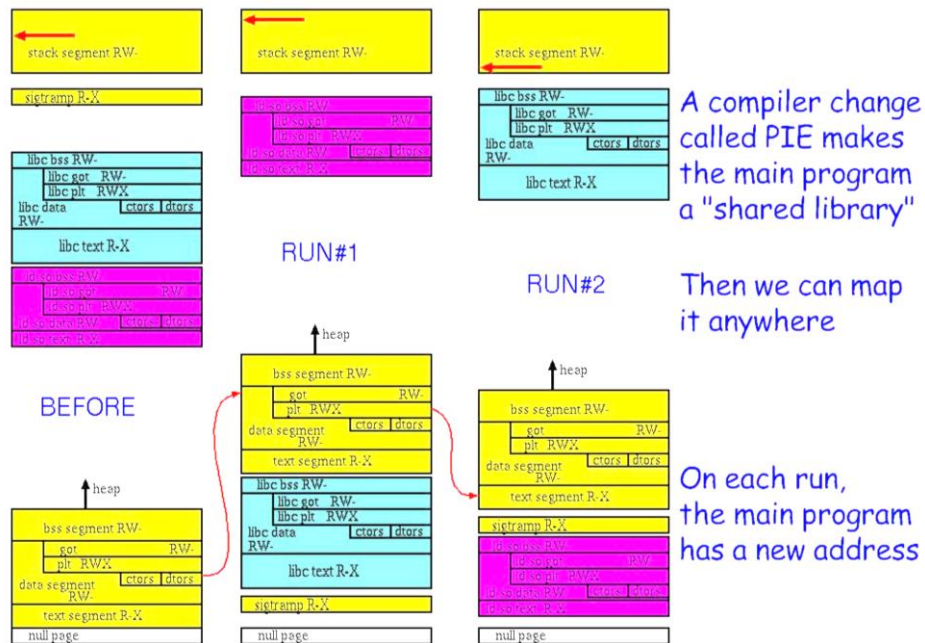


# START WITH LIBRARIES

ASLR: randomly map & order libraries



## PIE - Position Independent Executable



**ADD EXECUTABLES**





# FINALLY, DYNAMIC ALLOCATIONS

The word 'mmap' is centered within a light pink rounded rectangle. This rectangle is partially overlapped by a solid orange rounded rectangle on its top-left side. The entire composition is enclosed within a thin orange rounded rectangular border.

mmap

The word 'malloc' is centered within a light pink rounded rectangle. This rectangle is partially overlapped by a solid orange rounded rectangle on its top-left side. The entire composition is enclosed within a thin orange rounded rectangular border.

malloc



# LIMITATIONS OF ASLR

1. **Boot-time based randomization**
2. **Unsupported executables/libraries, low-entropy.**
3. **ASLR does not *trap* the attack**
4. **ASLR does not alert in a case of an attack**
5. **ASLR does not *provide information* about an attack**
6. **ASLR is being bypassed by exploits daily**

Posted by **MORDECHAI GURI, PH.D.** on December 17, 2015



# MAKING VIOLATIONS LESS DANGEROUS


**W^X  
Permissions**

**rodata**



# W | X PERMISSIONS

Many bugs are exploitable because the address space has memory that is both writeable and executable (permissions = W | X)



this location has to be executable for the exploit to work

We could make the stack non-executable...

Hmmmm... how about a generic policy for the whole address space:

A page may be either writeable or executable, but not both (unless the program specifically requests)

We call this policy  $W \wedge X$  ( $W \text{ xor } X$ )



# THE .RODATA SEGMENT

## W^X Transition: The .rodata segment

Readonly strings and pointers were stored in the `.text` segment: X | R

Meaning const data could be executed (could be code an attacker could use as ROP payload)

Solution: start using the ELF `.rodata` segment


These objects are now only R, lost their X permission

Greater policy: "minimal set of permissions"



# FINALLY, BLOCKING EXPLOITS

## Stack Protector



A typical stack frame...

Random value is inserted here by function prologue ...  
... and checked by function epilogue

Reordering: Arrays (strings) placed closer to random value -- integers and pointers placed further away

`-fstack-protector-all` compiled system is 1.3% slower at make build



# return-Oriented PROGRAMMING

**David Brumley**

Carnegie Mellon University



Credit: Some slides from Ed Schwartz

# ROP OVERVIEW

## **Idea:**

We forge shell code out of existing application logic gadgets

## **Requirements:**

vulnerability + gadgets + some unrandomized or predictable code  
(we need to know the addresses of gadgets)




# MOTIVATION: RETURN-TO-LIBC ATTACK

ret transfers control to  
system, which finds  
arguments on stack

Overwrite return address with address of libc  
“system” function

- setup fake return address and argument(s)
- ret will “call” libc function

**No injected code!**




# QUESTION

Randomized!

With ASLR, we cannot forge a correct value for ptr since ASLR will randomize addresses.

**What can we do?**

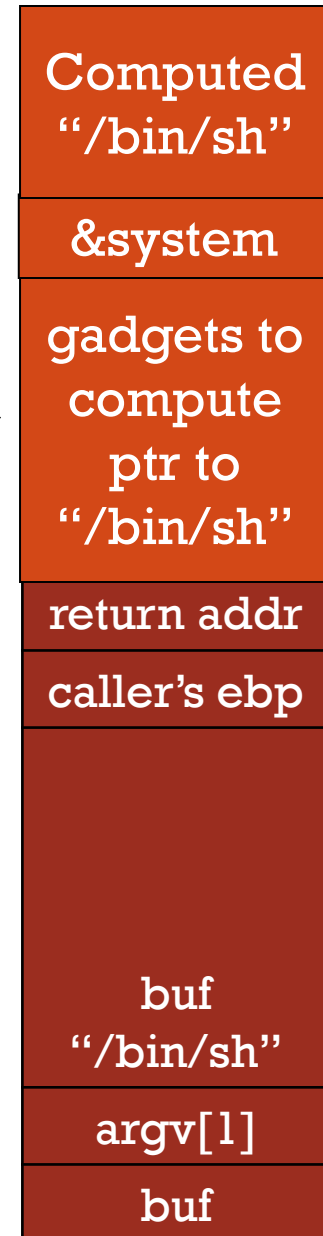



## Idea!

Get a copy of ESP to calculate address of “/bin/sh” on randomized stack.

This works because ASLR only protects against knowing *absolute* addresses, while we will find it's *relative address*.

**Writes**



# RETURN CHAINING

Suppose we want to call 2 functions in our exploit:

**foo**(arg1, arg2)

**bar**(arg3, arg4)

- Stack unwinds up
- First function returns into code to advance stack pointer
  - e.g., pop; pop; ret

What does this do?

Overwritten ret addr

arg4
arg3
&(pop-pop-ret)
<b>bar</b>
arg2
arg1
&(pop-pop-ret)
<b>foo</b>

# RETURN CHAINING

- When **foo** is executing, &pop-pop-ret is at the saved EIP slot.
- When **foo** returns, it executes pop-pop-ret to clear up arg1 (pop), arg2 (pop), and transfer control to **bar** (ret)

arg4
arg3
&(pop-pop-ret)
<b>bar</b>
arg2
arg1
&(pop-pop-ret)
<b>foo</b>




# THERE ARE MANY *SEMANTICALLY* *EQUIVALENT* WAYS TO ACHIEVE THE SAME NET SHELLCODE EFFECT

Let's practice thinking in gadgets

# AN EXAMPLE OPERATION

**Mem[v2] = v1**

**Desired Logic**



$a_1$ : mov eax, [esp] ; eax has v1

$a_2$ : mov ebx, [esp+8] ; ebx has v2

$a_3$ : mov [ebx], eax ; Mem[v2] = eax


**Implementation 1**

# IMPLEMENTING WITH ASSEMBLY

**Mem[v2] = v1**

**Desired Logic**

eax	<b>v<sub>1</sub></b>
ebx	
eip	a <sub>1</sub>



**Stack**

a<sub>1</sub>: pop eax  
a<sub>2</sub>: ret  
a<sub>3</sub>: pop ebx  
a<sub>4</sub>: ret  
a<sub>5</sub>: mov [ebx], eax

**Implementation 2**




# IMPLEMENTING WITH ASSEMBLY

**Mem[v2] = v1**

**Desired Logic**

eax	v <sub>1</sub>
ebx	
eip	a <sub>3</sub>



**Stack**

a<sub>1</sub>: pop eax  
a<sub>2</sub>: ret  
a<sub>3</sub>: pop ebx  
a<sub>4</sub>: ret  
a<sub>5</sub>: mov [ebx], eax


**Implementation 2**

# IMPLEMENTING WITH ASSEMBLY

**Mem[v2] = v1**

**Desired Logic**

eax	v <sub>1</sub>
ebx	v <sub>2</sub>
eip	a <sub>3</sub>



**Stack**

```
a1: pop  eax
a2: ret
a3: pop  ebx
a4: ret
a5: mov  [ebx], eax
```


**Implementation 2**

# IMPLEMENTING WITH ASSEMBLY

**Mem[v2] = v1**

**Desired Logic**

eax	v <sub>1</sub>
ebx	v <sub>2</sub>
eip	a <sub>5</sub>



**Stack**

```
a1: pop eax;  
a2: ret  
a3: pop ebx;  
a4: ret  
a5: mov [ebx], eax
```


**Implementation 2**

# IMPLEMENTING WITH GADGETS!

**Mem[v2] = v1**

**Desired Logic**

eax	v <sub>1</sub>
ebx	v <sub>2</sub>
eip	a <sub>5</sub>



**Stack**

```
a1: pop  eax;
a2: ret
a3: pop  ebx;
a4: ret
a5: mov  [ebx], eax
```

**Gadget 1** (points to a<sub>1</sub> and a<sub>2</sub>)  
**Gadget 2** (points to a<sub>3</sub> and a<sub>4</sub>)

**Implementation 2**

# EQUIVALENCE


**Mem[v2] = v1**

**Desired Logic**

semantically  
equivalent

a<sub>1</sub>: mov eax, [esp]  
a<sub>2</sub>: mov ebx, [esp+8]  
a<sub>3</sub>: mov [ebx], eax

**Implementation 1**



**Stack**

“Gadgets”

a<sub>1</sub>: pop eax; ret  
a<sub>2</sub>: pop ebx; ret  
a<sub>3</sub>: mov [ebx], eax

**Implementation 2**

# GADGETS

- A gadget is a set of instructions for carrying out a semantic action
  - mov, add, etc.
- Gadgets typically have a number of instructions
  - One instruction = native instruction set
  - More instructions = synthesize <- ROP
- Gadgets in ROP generally (but not always) end in return


# RETURN-ORIENTED PROGRAMMING (ROP)

**Mem[v2] = v1**

**Desired *Shellcode***

a<sub>1</sub>: pop eax; ret  
a<sub>2</sub>: pop ebx; ret  
a<sub>3</sub>: mov [ebx], eax

**Desired store executed!**



# Return-Oriented Programming

is A lot like a ransom  
note, BUT instead of cutting  
cut letters from magazines,  
YOU ARE cutting out  
instructions from text  
segments



# RO(P?) PROGRAMMING

1. Disassemble code
2. Identify useful code sequences as gadgets
3. Assemble gadgets into desired shellcode